

EXP 3: Map Reduce program to process a weather dataset.

AIM:

To implement MapReduce program to process a weather dataset.

Procedure:

Step 1: Create Data File:

Create a file named "word_count_data.txt" and populate it with text data that you wish to analyse.

Login with your hadoop user.

Download the dataset (weather data)

Output:

Step 2: Mapper Logic - mapper.py:

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

```
nano mapper.py
# Copy and paste the mapper.py code

#!/usr/bin/env python

import sys

# input comes from STDIN (standard input)
# the mapper will get daily max temperature and group it by month. so output will be
(month,daily_max_temperature)

for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # See the README hosted on the weather website which help us understand how each
    position represents a column    month = line[10:12]    daily_max = line[38:45]
    daily_max = daily_max.strip()
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be go through the shuffle process and then
        # be the input for the Reduce step, i.e. the input for reducer.py
```

```
#
# tab-delimited; month and daily max temperature as output
print ('%s\t%s' % (month ,daily_max))
```

Step 3: Reducer Logic - reducer.py:

Create a file named "reducer.py" to implement the logic for the reducer. The reducer will aggregate the occurrences of each word and generate the final output.

```
nano reducer.py
# Copy and paste the reducer.py code
```

reducer.py

```
#!/usr/bin/env python

from operator import itemgetter
import sys

#reducer will get the input from stdid which will be a collection of key, value(Key=month ,
value= daily max temperature)
#reducer logic: will get all the daily max temperature for a month and find max temperature
for the month
#shuffle will ensure that key are sorted(month)
current_month = None
current_max = 0
month = None

# input comes from STDIN for
line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # parse the input we got from mapper.py
    month, daily_max = line.split('\t', 1)

    # convert daily_max (currently a string) to float
    try:
        daily_max = float(daily_max)
    except ValueError:
        # daily_max was not a number, so silently
        # ignore/discard this line
        continue

    # this IF-switch only works because Hadoop shuffle process sorts map output
    # by key (here: month) before it is passed to the reducer
    if current_month == month:
        if daily_max >
```

```

current_max:        current_max = daily_max    else:
if current_month:
    # write result to STDOUT
    print ('%s\t%s' % (current_month, current_max))
current_max = daily_max
current_month = month

# output of the last month if current_month ==
month:    print ('%s\t%s' % (current_month,
current_max))

```

Step 4: Prepare Hadoop Environment:

Start the Hadoop daemons and create a directory in HDFS to store your data.

```
start-all.sh
```

Step 6: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.

```
chmod 777 mapper.py reducer.py
```

Step 7: Run the program using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access.

Then run the program using Hadoop Streaming.

```
hadoop fs -mkdir -p /weatherdata
```

```
hadoop fs -copyFromLocal /home/sx/Downloads/dataset.txt /weatherdata
```

```
hdfs dfs -ls /weatherdata
```

```

hadoop jar /home/sx/hadoop-3.2.3/share/hadoop/tools/lib/hadoop-streaming-3.2.3.jar \
-input /weatherdata/dataset.txt \
-output /weatherdata/output \
-file "/home/sx/Downloads/mapper.py" \
-mapper "python3 mapper.py" \
-file "/home/sx/Downloads/reducer.py" \
-reducer "python3 reducer.py"

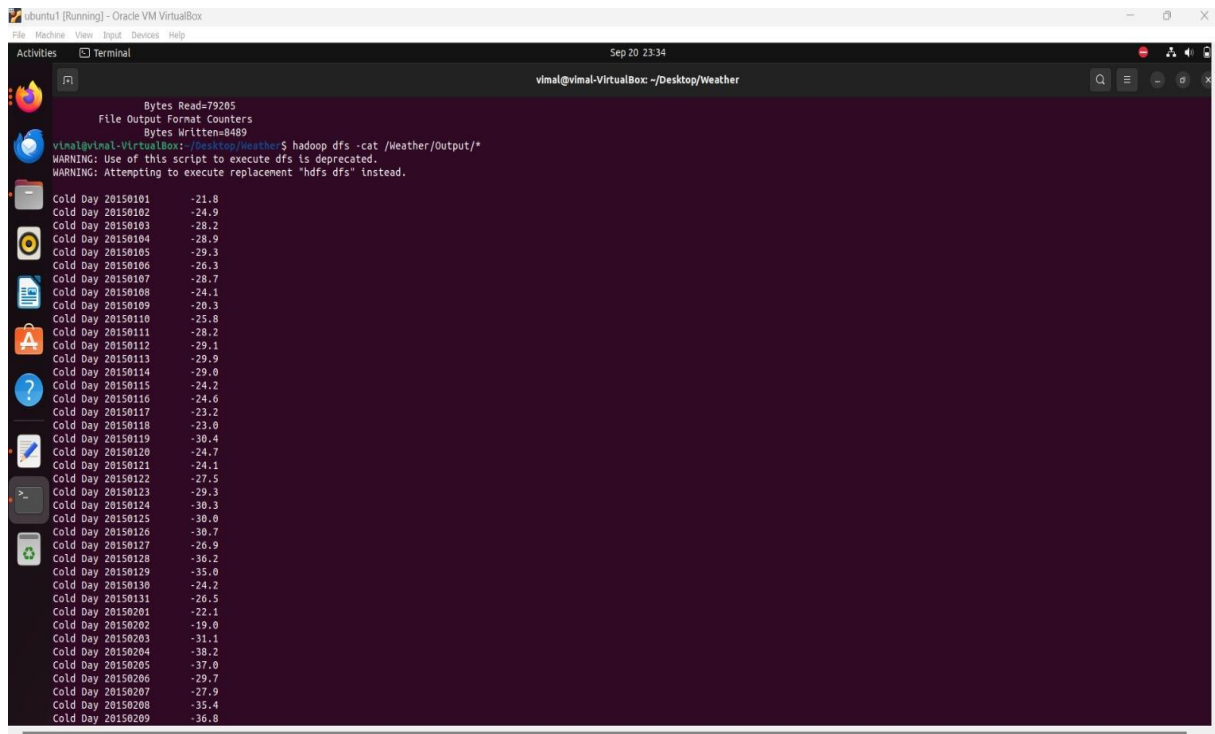
```

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/outputfile.txt
```

Step 8: Check Output:

Check the output of the program in the specified HDFS output directory.

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/output/ /part-00000
```



```
ubuntu1 [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Sep 20 23:34
vimal@vimal-VirtualBox: ~/Desktop/Weather

Bytes Read:79265
File Output Format Counters
Bytes Written:8489
vimal@vimal-VirtualBox:~/Desktop/Weather$ hadoop dfs -cat /Weather/Output/*
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement 'hdfs dfs' instead.
Cold Day 20150101 -21.8
Cold Day 20150102 -24.9
Cold Day 20150103 -28.2
Cold Day 20150104 -28.9
Cold Day 20150105 -29.3
Cold Day 20150106 -26.3
Cold Day 20150107 -28.7
Cold Day 20150108 -24.1
Cold Day 20150109 -20.3
Cold Day 20150110 -25.8
Cold Day 20150111 -28.2
Cold Day 20150112 -29.1
Cold Day 20150113 -29.9
Cold Day 20150114 -29.0
Cold Day 20150115 -24.2
Cold Day 20150116 -24.6
Cold Day 20150117 -23.2
Cold Day 20150118 -23.0
Cold Day 20150119 -30.4
Cold Day 20150120 -24.7
Cold Day 20150121 -24.1
Cold Day 20150122 -27.5
Cold Day 20150123 -29.3
Cold Day 20150124 -30.3
Cold Day 20150125 -30.0
Cold Day 20150126 -30.7
Cold Day 20150127 -26.9
Cold Day 20150128 -36.2
Cold Day 20150129 -35.0
Cold Day 20150130 -24.2
Cold Day 20150131 -26.5
Cold Day 20150201 -22.1
Cold Day 20150202 -19.0
Cold Day 20150203 -31.1
Cold Day 20150204 -38.2
Cold Day 20150205 -37.0
Cold Day 20150206 -29.7
Cold Day 20150207 -27.9
Cold Day 20150208 -35.4
Cold Day 20150209 -36.8
```

After copy and paste the above output in your local file give the below command to remove the directory from hdfs : `hadoop fs -rm -r /weatherdata/output`

Result:

Thus, the program for weather dataset using Map Reduce has been executed successfully.